

# Assignment 24 Data Science Masters

February 18, 2019

## 0.1 Session 24 -Assignment - Machine Learning 5

### 0.1.1 House Pricing from Boston Data Set.

#### Load Libraries

```
In [24]: # Core Libraries to load (for data manipulation and analysis)
import pandas as pd
import numpy as np
import math
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# Import stats module from scipy for Statistical analysis
import scipy.stats as stats

# Core Libraries to load - Machine Learning
import sklearn

## Import LinearRegression Module - Modelling
from sklearn.linear_model import LinearRegression

## Import RandomForestRegressor Module - Modelling
from sklearn.ensemble import RandomForestRegressor

## Import StandardScaler - Model Data Preprocessing
from sklearn.preprocessing import StandardScaler

## Import train_test_split Module
from sklearn.model_selection import train_test_split, GridSearchCV, ShuffleSplit, cross_val_score

## Importing mean_squared_error and r2_score from sklearn.metrics
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

## Import boston dataset from sklearn.datasets
from sklearn.datasets import load_boston
```

```

# Warnings Library - Ignore warnings
import warnings
warnings.filterwarnings('ignore')

### Load Boston dataset into a variable
boston = load_boston()

```

## 0.1.2 Understand the Dataset and the Data

```

In [3]: # Get keys of boston dataset dictionary
print(boston.keys())

```

```
dict_keys(['data', 'target', 'feature_names', 'DESCR'])
```

```

In [4]: # Get the number of rows and columns in the dataset
boston.data.shape

```

```
Out[4]: (506, 13)
```

```

In [5]: # Get the column names in the dataset
print(boston.feature_names)

```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

```

In [6]: # Get description of the column names in the dataset
print(boston.DESCR)

```

```

Boston House Prices dataset
=====

```

Notes

-----

Data Set Characteristics:

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive

:Median Value (attribute 14) is usually the target

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)

- NOX        nitric oxides concentration (parts per 10 million)
- RM        average number of rooms per dwelling
- AGE        proportion of owner-occupied units built prior to 1940
- DIS        weighted distances to five Boston employment centres
- RAD        index of accessibility to radial highways
- TAX        full-value property-tax rate per \$10,000
- PTRATIO    pupil-teacher ratio by town
- B         $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town
- LSTAT     % lower status of the population
- MEDV     Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<http://archive.ics.uci.edu/ml/datasets/Housing>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

#### **\*\*References\*\***

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity'
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the AAAI Conference on Artificial Intelligence
- many more! (see <http://archive.ics.uci.edu/ml/datasets/Housing>)

```
In [7]: # Create dataframe from boston.data and target information in variables
        features = pd.DataFrame(boston.data, columns=boston.feature_names)
        targets = boston.target
```

```
In [8]: # Updating the dataframe by adding the target column
        features["PRICE"] = targets
```

#### ***Information of the dataframe created from the data set***

```
In [9]: # Shape of the dataframe
        features.shape
```

```
Out[9]: (506, 14)
```

```
In [10]: # Dataframe after updating with target
features.head(5)
```

```
Out[10]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
In [11]: features.get_dtype_counts()
```

```
Out[11]: float64    14
dtype: int64
```

*The columns' datatypes are all numeric*

## 0.2 Basic Details about Data - For Data Cleaning and Data Wrangling

*Check for datatypes and presence of null values using info()*

```
In [12]: features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
CRIM      506 non-null float64
ZN        506 non-null float64
INDUS     506 non-null float64
CHAS      506 non-null float64
NOX       506 non-null float64
RM        506 non-null float64
AGE       506 non-null float64
DIS       506 non-null float64
RAD       506 non-null float64
TAX       506 non-null float64
PTRATIO   506 non-null float64
B         506 non-null float64
LSTAT     506 non-null float64
PRICE     506 non-null float64
```

```
dtypes: float64(14)
memory usage: 55.4 KB
```

*No cleaning required as the data is already cleaned and has no null or NaN values*

*Checking if there are any row values = zero that need our consideration so that we can decide to study those rows*

```
In [15]: # Checking for the number of rows containing all values = 0
features.loc[(features==0).all(axis=1)].shape
```

```
Out[15]: (0, 14)
```

```
In [16]: # Checking for the number of rows containing atleast one value = 0
features.loc[(features==0).any(axis=1)].shape
```

```
Out[16]: (499, 14)
```

*The CHAS column is River dummy variable(= 1 if tract bounds river; 0 otherwise) and therefore the zeros in the column are valid values. So we don't need to clean the data in that column*

```
In [19]: features.columns.values
```

```
Out[19]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
               'TAX', 'PTRATIO', 'B', 'LSTAT', 'PRICE'], dtype=object)
```

## 1 Basic Statistical Information

```
In [20]: # Getting basic statistical information about the columns
features.describe() # Only numerical columns
```

```
Out[20]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.593761	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.596783	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.647423	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	AGE	DIS	RAD	TAX	PTRATIO	B	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000	

max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000
-----	------------	-----------	-----------	------------	-----------	------------

	LSTAT	PRICE
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
In [21]: # Getting correlation between various numerical columns
features.corr()
```

```
Out [21]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
CRIM	1.000000	-0.199458	0.404471	-0.055295	0.417521	-0.219940	0.350784	
ZN	-0.199458	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	
INDUS	0.404471	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	
CHAS	-0.055295	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	
NOX	0.417521	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	
RM	-0.219940	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	
AGE	0.350784	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	
DIS	-0.377904	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	
RAD	0.622029	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	
TAX	0.579564	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	
PTRATIO	0.288250	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	
B	-0.377365	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	
LSTAT	0.452220	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	
PRICE	-0.385832	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	

	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
CRIM	-0.377904	0.622029	0.579564	0.288250	-0.377365	0.452220	-0.385832
ZN	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
LSTAT	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
PRICE	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000

*The column RM is highly correlated with PRICE. (Correlation = 0.695)*

*The column LSTAT is highly negatively correlated with PRICE. (Correlation = -0.737)*  
*The columns INDUS, NOX, TAX, PTRATIO are moderately negatively correlated with PRICE (-0.7 < Correlations < -0.4)*

## 2 Data Exploration - Visual Analysis OR Exploratory Analysis

### 2.1 Uni-variate

In [22]: # Plotting the histograms of numerical columns to understand their distribution  
 features.hist(bins=50, figsize=(20,20), layout=(5,3))  
 plt.show()




## 2.2 Bi-Variate

*Plotting: PRICE vs RM, LSTAT, INDUS, NOX, TAX and PTRATIO*

```
In [25]: sns.jointplot(x=features["RM"], y=features["PRICE"], kind='scatter',size = 8)
```


```
Out[25]: <seaborn.axisgrid.JointGrid at 0x29ba6f56b00>
```



```
In [26]: sns.jointplot(x=features["LSTAT"], y=features["PRICE"], kind='scatter',size = 8)
```


```
Out[26]: <seaborn.axisgrid.JointGrid at 0x29ba73b4470>
```






```
In [27]: sns.jointplot(x=features["INDUS"], y=features["PRICE"], kind='scatter',size = 8)
```

```
Out[27]: <seaborn.axisgrid.JointGrid at 0x29ba7503780>
```




```
In [28]: sns.jointplot(x=features["NOX"], y=features["PRICE"], kind='scatter',size = 8)
```

```
Out[28]: <seaborn.axisgrid.JointGrid at 0x29ba6122be0>
```




```
In [29]: sns.jointplot(x=features["TAX"], y=features["PRICE"], kind='scatter',size = 8)
```

```
Out[29]: <seaborn.axisgrid.JointGrid at 0x29ba67762b0>
```




```
In [30]: sns.jointplot(x=features["INDUS"], y=features["PRICE"], kind='scatter',size = 8)
```

```
Out[30]: <seaborn.axisgrid.JointGrid at 0x29ba69225c0>
```



```
In [31]: sns.jointplot(x=features["INDUS"], y=features["PTRATIO"], kind='scatter',size = 8)
```

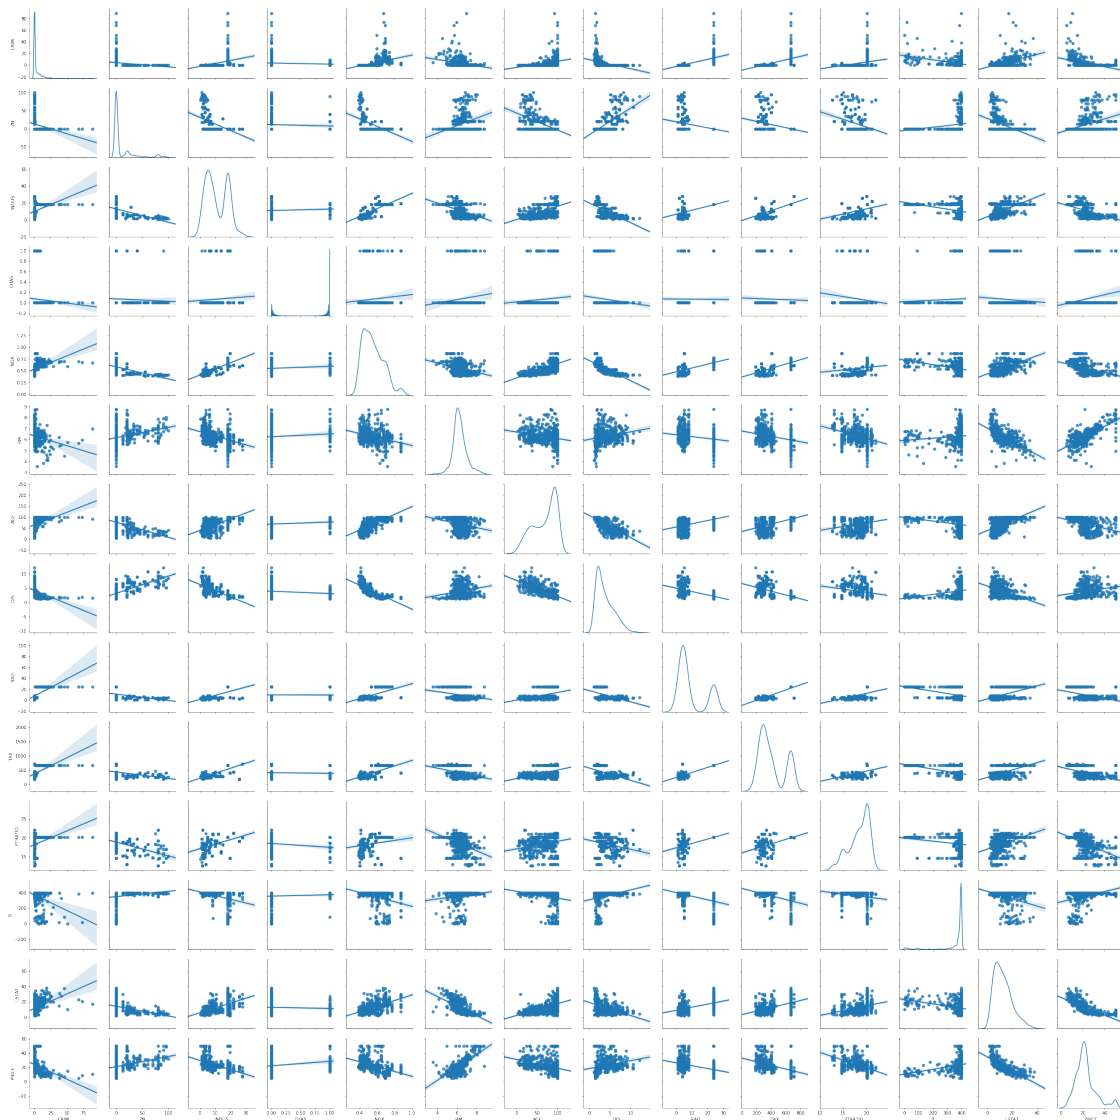
```
Out[31]: <seaborn.axisgrid.JointGrid at 0x29ba6d0ce10>
```



## 2.3 Multi-Variate


In [32]: `sns.pairplot(features, kind='reg', diag_kind = 'kde')`

Out [32]: `<seaborn.axisgrid.PairGrid at 0x29ba6b41ac8>`



```
In [33]: # Checking for correlations using HEATMAP
plt.figure(figsize=(20,20))
sns.heatmap(features.corr(), cmap="PRGn", annot= True)
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x29baf130c50>
```



### 3 Feature Engineering

*There is no need to engineer features from this dataset. Also, there are no categorical variables to engineer*

### 4 Train - Test Split

```
In [34]: X = features.drop("PRICE", axis=1)
         Y = features["PRICE"]
```



```

# We will be using 80:20 split for train and test datasets
x_train, x_test, y_train, y_test = train_test_split(X,Y,test_size=0.20, random_state =

In [35]: print(x_train.shape, x_test.shape)

(404, 13) (102, 13)

In [36]: print(y_train.shape, y_test.shape)

(404,) (102,)

```

## 5 Fitting Models

### 5.1 Linear Regression

```

In [38]: lm = LinearRegression()
         model = lm.fit(x_train, y_train) # Sklearn already considers the intercepts for linear
         print("Estimated Beta Coefficients: \n", model.coef_)

         y_test_pred = model.predict(x_test)

         print("\nLinear Regression - Base", "\n\t R2-Score:", model.score(x_test, y_test),
               "\n\t RMSE:", math.sqrt(mean_squared_error(y_test_pred, y_test)))

Estimated Beta Coefficients:
[-8.01644009e-02  4.79926054e-02 -5.07131765e-03  3.06486600e+00
 -1.61596810e+01  3.66858142e+00 -8.46805789e-03 -1.51719956e+00
  2.86612524e-01 -1.21155515e-02 -9.24761912e-01  9.62688265e-03
 -4.86676845e-01]

Linear Regression - Base
R2-Score: 0.7554467329645207
RMSE: 4.860294126345348

```

### 5.2 RandomForestRegressor

```

In [39]: rf_reg = RandomForestRegressor(n_estimators=100)
         rf_model= rf_reg.fit(x_train, y_train)
         y_test_pred = rf_model.predict(x_test)

         print("## RandomForest Regressor - Unscaled Data", "\n\t R2-Score:", rf_model.score(x_test, y_test),
               "\n\t RMSE:", math.sqrt(mean_squared_error(y_test_pred, y_test)))

## RandomForest Regressor - Unscaled Data
R2-Score: 0.8912837096319373

```

RMSE: 3.2405828542554054

### 5.3 RandomForestRegressor with StandardScaled data

```
In [40]: scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)
scaled_features = pd.DataFrame(scaled_features, columns = features.columns.values)
scaled_features.head()
```

```
Out [40]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
0	-0.417713	0.284830	-1.287909	-0.272599	-0.144217	0.413672	-0.120013	
1	-0.415269	-0.487722	-0.593381	-0.272599	-0.740262	0.194274	0.367166	
2	-0.415272	-0.487722	-0.593381	-0.272599	-0.740262	1.282714	-0.265812	
3	-0.414680	-0.487722	-1.306878	-0.272599	-0.835284	1.016303	-0.809889	
4	-0.410409	-0.487722	-1.306878	-0.272599	-0.835284	1.228577	-0.511180	

	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.140214	-0.982843	-0.666608	-1.459000	0.441052	-1.075562	0.159686
1	0.557160	-0.867883	-0.987329	-0.303094	0.441052	-0.492439	-0.101524
2	0.557160	-0.867883	-0.987329	-0.303094	0.396427	-1.208727	1.324247
3	1.077737	-0.752922	-1.106115	0.113032	0.416163	-1.361517	1.182758
4	1.077737	-0.752922	-1.106115	0.113032	0.441052	-1.026501	1.487503

```
In [41]: X_scaled = scaled_features.drop("PRICE", axis=1)
Y_scaled = scaled_features["PRICE"]

# We will be using 80:20 split for train and test datasets
x_scaled_train, x_scaled_test, y_scaled_train, y_scaled_test = train_test_split(X_scaled, Y_scaled,
                                        test_size=0.2, random_state=42)

In [43]: rf_reg = RandomForestRegressor(n_estimators=100)
rf_model_scaled = rf_reg.fit(x_scaled_train, y_scaled_train)
y_scaled_test_pred = rf_model_scaled.predict(x_scaled_test)

print("## RandomForestRegressor - ScaledData", "\n\t R2-Score:", rf_model_scaled.score(x_scaled_test, y_scaled_test),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_scaled_test, y_scaled_test_pred)))


## RandomForestRegressor - ScaledData
R2-Score: 0.8843132183938184
RMSE: 0.36382802988641866
```

## 6 Feature Selection

```
In [44]: importance = pd.DataFrame.from_dict({'cols':x_train.columns, 'importance': rf_reg.feature_importances_})
importance = importance.sort_values(by='importance', ascending=False)
```

```
plt.figure(figsize=(20,15))
sns.barplot(importance.cols, importance.importance)
plt.xticks(rotation=90)
```

Out[44]: (array([ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]),  
<a list of 13 Text xticklabel objects>)



```
In [46]: imp_cols = importance[importance.importance >= 0.01].cols.values
         imp_cols
```

Out[46]: array(['RM', 'LSTAT', 'DIS', 'CRIM', 'TAX', 'PTRATIO', 'AGE', 'NOX', 'B'],  
dtype=object)

```
In [47]: # Fitting models with columns where feature importance>=0.01
```

```
x_train, x_test, y_train, y_test = train_test_split(X[imp_cols],Y,test_size=0.20, ran

rf_reg = RandomForestRegressor(n_estimators=100)
rf_model= rf_reg.fit(x_train, y_train)
y_test_pred = rf_model.predict(x_test)
```

```

print("## RandomForest Regressor - Unscaled Data", "\n\t R2-Score:", rf_model.score(x_train, y_train),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_test_pred, y_test)))

x_scaled_train, x_scaled_test, y_scaled_train, y_scaled_test = train_test_split(X_scaled, y_scaled,
                                          test_size=0.2, random_state=42)

rf_reg = RandomForestRegressor(n_estimators=100)
rf_model_scaled = rf_reg.fit(x_scaled_train, y_scaled_train)
y_scaled_test_pred = rf_model_scaled.predict(x_scaled_test)

print("## RandomForestRegressor - Scaled Data", "\n\t R2-Score:", rf_model_scaled.score(x_scaled_train, y_scaled_train),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_scaled_test_pred, y_scaled_test)))

## RandomForest Regressor - Unscaled Data
R2-Score: 0.8883031363395156
RMSE: 3.2847045096456196

## RandomForestRegressor - Scaled Data
R2-Score: 0.8933982831642273
RMSE: 0.34924997580717

In [48]: imp_cols = importance[importance.importance >= 0.005].cols.values
         imp_cols

Out[48]: array(['RM', 'LSTAT', 'DIS', 'CRIM', 'TAX', 'PTRATIO', 'AGE', 'NOX', 'B',
                'INDUS'], dtype=object)

In [49]: # Fitting models with columns where feature importance>=0.005

x_train, x_test, y_train, y_test = train_test_split(X[imp_cols], Y, test_size=0.2, random_state=42)

rf_reg = RandomForestRegressor(n_estimators=100)
rf_model = rf_reg.fit(x_train, y_train)
y_test_pred = rf_model.predict(x_test)

print("## RandomForest Regressor - Unscaled Data", "\n\t R2-Score:", rf_model.score(x_train, y_train),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_test_pred, y_test)))

x_scaled_train, x_scaled_test, y_scaled_train, y_scaled_test = train_test_split(X_scaled, y_scaled,
                                          test_size=0.2, random_state=42)

rf_reg = RandomForestRegressor(n_estimators=100)
rf_model_scaled = rf_reg.fit(x_scaled_train, y_scaled_train)
y_scaled_test_pred = rf_model_scaled.predict(x_scaled_test)

```

```

print("## RandomForestRegressor - Scaled Data", "\n\t R2-Score:", rf_model_scaled.score(x_scaled_test, y_scaled_test_pred),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_scaled_test_pred, y_scaled_test)))

## RandomForest Regressor - Unscaled Data
R2-Score: 0.888963184316594
RMSE: 3.274985011191563

## RandomForestRegressor - Scaled Data
R2-Score: 0.8840231988338807
RMSE: 0.36428379153493723

In [50]: imp_cols = importance[importance.importance >= 0.002].cols.values
         imp_cols

Out[50]: array(['RM', 'LSTAT', 'DIS', 'CRIM', 'TAX', 'PTRATIO', 'AGE', 'NOX', 'B',
               'INDUS', 'RAD'], dtype=object)

In [51]: # Fitting models with columns where feature importance>=0.002

x_train, x_test, y_train, y_test = train_test_split(X[imp_cols], Y, test_size=0.20, random_state=42)

rf_reg = RandomForestRegressor(n_estimators=100)
rf_model = rf_reg.fit(x_train, y_train)
y_test_pred = rf_model.predict(x_test)

print("## RandomForest Regressor - Unscaled Data", "\n\t R2-Score:", rf_model.score(x_test, y_test),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_test_pred, y_test)))

x_scaled_train, x_scaled_test, y_scaled_train, y_scaled_test = train_test_split(X_scaled, Y_scaled, test_size=0.20, random_state=42)

rf_reg = RandomForestRegressor(n_estimators=100)
rf_model_scaled = rf_reg.fit(x_scaled_train, y_scaled_train)
y_scaled_test_pred = rf_model_scaled.predict(x_scaled_test)

print("## RandomForestRegressor - Scaled Data", "\n\t R2-Score:", rf_model_scaled.score(x_scaled_test, y_scaled_test_pred),
      "\n\t RMSE:", math.sqrt(mean_squared_error(y_scaled_test_pred, y_scaled_test)))

## RandomForest Regressor - Unscaled Data
R2-Score: 0.8811702177745246
RMSE: 3.3879615118655564

## RandomForestRegressor - Scaled Data
R2-Score: 0.8888968021878442
RMSE: 0.35654763879104023

```

## 7 Validation

```
In [52]: # Cross validating the model created with columns whose feature importances >= 0.002,
scoring = 'neg_mean_squared_error'
kfold = KFold(n_splits=10, random_state=100)

cv_results = cross_val_score(model, x_train,y_train, cv=kfold, scoring=scoring)
print("## Linear Regression","\n\t CV-Mean:", cv_results.mean(),
      "\n\t CV-Std. Dev:", cv_results.std())

cv_results = cross_val_score(rf_model, x_train,y_train, cv=kfold, scoring=scoring)
print("## RandomForestRegressor - Unscaled data","\n\t CV-Mean:", cv_results.mean(),
      "\n\t CV-Std. Dev:", cv_results.std())

cv_results = cross_val_score(rf_model_scaled, x_scaled_train,y_scaled_train, cv=kfold)
print("## RandomForestRegressor - Scaled data","\n\t CV-Mean:", cv_results.mean(),
      "\n\t CV-Std. Dev:", cv_results.std())

## Linear Regression
CV-Mean: -24.501361520600785
CV-Std. Dev: 7.478252464775003

## RandomForestRegressor - Unscaled data
CV-Mean: -10.768933929329275
CV-Std. Dev: 4.507874541914067

## RandomForestRegressor - Scaled data
CV-Mean: -0.13622902633058565
CV-Std. Dev: 0.060291200294718905
```

## 8 Optimization - Model Tuning

```
In [53]: RF_Regressor = RandomForestRegressor(n_estimators=100, n_jobs = -1, random_state = 100)

CV = ShuffleSplit(test_size=0.20, random_state=100)

param_grid = {"max_depth": [5, None],
              "n_estimators": [50, 100, 150, 200],
              "min_samples_split": [2, 4, 5],
              "min_samples_leaf": [2, 4, 6]
              }
```

### 8.1 Best Estimator - Unscaled Data

```
In [54]: rscv_grid1 = GridSearchCV(RF_Regressor, param_grid=param_grid, verbose=1)
```

```
In [55]: rscv_grid1.fit(x_train, y_train)
```

Fitting 3 folds for each of 72 candidates, totalling 216 fits

```
[Parallel(n_jobs=1)]: Done 216 out of 216 | elapsed: 1.6min finished
```

```
Out [55]: GridSearchCV(cv=None, error_score='raise',
                      estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                      max_features='auto', max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
                      oob_score=False, random_state=100, verbose=0, warm_start=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'max_depth': [5, None], 'n_estimators': [50, 100, 150, 200], 'min_
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=1)
```

```
In [56]: rscv_grid1.best_params_
```

```
Out [56]: {'max_depth': None,
           'min_samples_leaf': 2,
           'min_samples_split': 2,
           'n_estimators': 50}
```

```
In [57]: # Best Estimator - Unscaled
rf_model = rscv_grid1.best_estimator_
rf_model.fit(x_train, y_train)
```

```
Out [57]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=2, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=50, n_jobs=-1,
                                oob_score=False, random_state=100, verbose=0, warm_start=False)
```

```
In [58]: rf_model.score(x_test, y_test)
```

```
Out [58]: 0.8781370746758193
```

## 8.2 Best Estimator - Scaled Data

```
In [59]: rscv_grid2 = GridSearchCV(RF_Regressor, param_grid=param_grid, verbose=1)
```

```
In [60]: rscv_grid2.fit(x_scaled_train, y_scaled_train)
```

Fitting 3 folds for each of 72 candidates, totalling 216 fits

```
[Parallel(n_jobs=1)]: Done 216 out of 216 | elapsed: 1.6min finished
```

```
Out [60]: GridSearchCV(cv=None, error_score='raise',
                      estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                      max_features='auto', max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=-1,
                      oob_score=False, random_state=100, verbose=0, warm_start=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'max_depth': [5, None], 'n_estimators': [50, 100, 150, 200], 'min_
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=1)
```

```
In [61]: rscv_grid2.best_params_
```

```
Out [61]: {'max_depth': None,
           'min_samples_leaf': 2,
           'min_samples_split': 4,
           'n_estimators': 150}
```

```
In [62]: # Best Estimator - Scaled
rf_model_scaled = rscv_grid2.best_estimator_
rf_model_scaled.fit(x_scaled_train, y_scaled_train)
```

```
Out [62]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=2, min_samples_split=4,
                                min_weight_fraction_leaf=0.0, n_estimators=150, n_jobs=-1,
                                oob_score=False, random_state=100, verbose=0, warm_start=False)
```

```
In [63]: rf_model_scaled.score(x_scaled_test, y_scaled_test)
```

```
Out [63]: 0.8771363772523068
```

## 9 Comparing Performance Metrics

```
In [64]: print("RandomForestRegressor - Unscaled Data\n\t R2-Score:", rf_model.score(x_test, y_test),
              "\n\t RMSE:", math.sqrt(mean_squared_error(rf_model.predict(x_test), y_test)))

print("RandomForestRegressor - Scaled Data\n\t R2-Score:", rf_model_scaled.score(x_scaled_test, y_scaled_test),
      "\n\t RMSE:", math.sqrt(mean_squared_error(rf_model_scaled.predict(x_scaled_test), y_scaled_test)))
```

```
RandomForestRegressor - Unscaled Data
R2-Score: 0.8781370746758193
RMSE: 3.430928100601086
```



```
RandomForestRegressor - Scaled Data  
R2-Score: 0.8771363772523068  
RMSE: 0.37494359878281436
```

## 10 Choosing the model

*We can see that Random Forest Regressor trained on scaled data gives better RMSE value (= 0.37447). So, Random Forest Regressor trained on scaled data should be used as the regression model for this dataset.*